

Maze Simulator and Reinforcement Learning Agent Report

Shivam Chaubey

February 9, 2025

1 Introduction

This report details the implementation of a maze simulator and a reinforcement learning (RL) agent designed to navigate the generated maze. The maze is generated using a randomized depth-first search algorithm. The RL agent employs Q-learning to find a path from a designated start point to an end point within the maze. The primary goal is to develop an agent that can find a reasonably short path through various maze configurations.

2 Code Overview

The Python code is structured into two main components: the `MazeSimulator` class and the `mazeSolverUsingRL` function.

2.1 MazeSimulator Class

The `MazeSimulator` class is responsible for generating and representing the maze.

2.1.1 `__init__` Method

This method initializes the maze with a specified width and height, creating a grid filled with walls (represented by 1). It also initializes the start and end points to `None`.

2.1.2 `generate_maze` Method

This method uses a depth-first search (DFS) algorithm to carve out paths in the maze. It starts from a random cell and recursively visits neighboring cells, marking them as part of the path (represented by 0). The `is_valid_move` helper function ensures that the DFS stays within the maze boundaries and only moves to unvisited cells.

2.1.3 `set_start_and_end` Method

This method randomly selects two distinct open cells (cells with value 0) in the maze and designates them as the start and end points for the agent.

2.1.4 `get_maze_array` Method

This method creates a 2D array representation of the maze, where 1 represents walls and 0 represents paths. It also marks the start and end points as 1 for representation purposes.

2.1.5 `save_maze_image` Method

This method uses Matplotlib to visualize the maze and save it as a PNG image. The start and end points are colored differently for clarity.

2.2 `mazeSolverUsingRL` Function

This function implements a Q-learning algorithm to train an agent to navigate the maze.

2.2.1 Q-Learning Implementation

The Q-learning algorithm uses a Q-table to store the expected rewards for taking specific actions in different states. The agent iteratively explores the maze, updating the Q-table based on the rewards it receives.

Hyperparameters:

- **alpha** (Learning rate): Determines the extent to which newly acquired information overrides old information.
- **gamma** (Discount factor): Determines the importance of future rewards.
- **epsilon** (Exploration rate): Controls the balance between exploration (trying new actions) and exploitation (using known optimal actions).
- **num_episodes**: The number of training iterations.

Q-Table Initialization: The Q-table is a 3D NumPy array that stores the Q-values for each state-action pair. The dimensions correspond to the height, width, and number of possible actions (up, right, down, left).

Action Mappings: The agent can take four actions: move up, right, down, or left. These actions are represented as coordinate changes: $(-1, 0)$, $(0, 1)$, $(1, 0)$, $(0, -1)$.

Reward Function: The reward function assigns rewards based on the agent's actions:

- Reaching the end point: +100
- Hitting a wall: -10
- Moving to an open cell: -1 (small penalty per step to encourage efficiency)

Q-Learning Loop: The Q-learning loop iterates through a specified number of episodes. In each episode, the agent starts from the start point and continues until it reaches the end point or hits a wall. The agent chooses an action based on an epsilon-greedy policy, which balances exploration and exploitation. The Q-table is updated using the Q-learning update rule.

2.2.2 Path Extraction

After the Q-table is trained, the function extracts the optimal path by starting at the start point and repeatedly choosing the action with the highest Q-value until it reaches the end point. A check is included to prevent the agent from getting stuck in loops.

3 Approach

The approach combines a standard maze generation algorithm (DFS) with a reinforcement learning technique (Q-learning) to solve the maze. This hybrid approach allows for the creation of complex mazes and the training of an agent to efficiently navigate them.

3.1 Maze Generation

The DFS algorithm ensures that the maze is fully connected, meaning there is a path from any cell to any other cell. The randomization aspect ensures that each generated maze is unique.

3.2 Reinforcement Learning

Q-learning is a model-free reinforcement learning algorithm that learns an optimal policy by iteratively updating a Q-table. The Q-table stores the expected rewards for taking specific actions in different states. The epsilon-greedy policy allows the agent to explore the maze initially and then gradually exploit its knowledge to find the shortest path.

4 Code Listing

```
1 import random
2 import matplotlib.pyplot as plt
3 import numpy as np
4
5 # Do not change the maze generation
6
7 class MazeSimulator:
8     def __init__(self, width, height):
9         self.width = width
10        self.height = height
11        self.maze = [[1 for _ in range(width)] for _ in range(height)] # 1 = wall, 0 =
path
12        self.start = None
13        self.end = None
14
15    def generate_maze(self):
16        def is_valid_move(x, y):
17            return 0 <= x < self.height and 0 <= y < self.width and self.maze[x][y] == 1
18
19        def dfs(x, y):
20            self.maze[x][y] = 0
21            directions = [(0, 2), (2, 0), (0, -2), (-2, 0)]
22            random.shuffle(directions)
23
24            for dx, dy in directions:
25                nx, ny = x + dx, y + dy
26                if is_valid_move(nx, ny):
27                    self.maze[x + dx // 2][y + dy // 2] = 0 # Remove the wall
28                    dfs(nx, ny)
29
30        start_x, start_y = random.randrange(0, self.height, 2), random.randrange(0, self
.width, 2)
31        dfs(start_x, start_y)
32
33    def set_start_and_end(self):
34        open_cells = [(x, y) for x in range(self.height) for y in range(self.width) if
self.maze[x][y] == 0]
35        self.start, self.end = random.sample(open_cells, 2)
36
37    def get_maze_array(self):
38        maze_array = [[1 if self.maze[x][y] == 0 or (x, y) in [self.start, self.end]
else 0 for y in range(self.width)] for x in range(self.height)]
39        return maze_array
40
41    def save_maze_image(self, filename="maze.png"):
42        maze_copy = [[2 if (x, y) == self.start else 3 if (x, y) == self.end else self.
maze[x][y] for y in range(self.width)] for x in range(self.height)]
43        plt.figure(figsize=(10, 10))
44        plt.imshow(maze_copy, cmap="viridis", origin="upper")
45        plt.axis("off")
46        plt.savefig(filename)
47        plt.close()
48
49
50    def mazeSolverUsingRL(self):
51        """
52        Solve the maze using Q-learning.
53
54        The RL agent learns a policy for moving from the start to the end cell.
55        Actions: 0 = Up, 1 = Right, 2 = Down, 3 = Left.
56        Rewards:
57        - +100 for reaching the end.
58        - -10 for hitting a wall (invalid move).
59        - -1 per step to encourage shorter paths.
60
61        Returns:
62        path (list of tuple): The sequence of (row, col) states representing the
optimal path.
63        """
64
65        # Hyperparameters
```

```

66     alpha = 0.1          # Learning rate
67     gamma = 0.9          # Discount factor
68     epsilon = 0.1        # Exploration rate
69     num_episodes = 100000 # Number of training episodes
70     max_steps = self.width * self.height * 4 # Maximum steps per episode (safeguard
71 )
72
73     # Initialize the Q-table with zeros.
74     # Dimensions: [height][width][number_of_actions]
75     q_table = np.zeros((self.height, self.width, 4))
76
77     # Define the four possible actions as (dx, dy) movements.
78     actions = [(-1, 0), # Up
79               (0, 1),   # Right
80               (1, 0),   # Down
81               (0, -1)]  # Left
82
83     def is_valid_state(x, y):
84         """Return True if (x, y) is within bounds and is an open cell (path)."""
85         return 0 <= x < self.height and 0 <= y < self.width and self.maze[x][y] == 0
86
87     def choose_action(state):
88         """Choose an action using an epsilon-greedy strategy."""
89         if random.uniform(0, 1) < epsilon:
90             return random.choice(range(4)) # Explore: choose a random action
91         else:
92             return np.argmax(q_table[state[0], state[1]]) # Exploit: choose the
93             best-known action
94
95     def get_reward(state):
96         """
97         Return the reward for moving into a given state.
98         - 100 if the state is the end.
99         - -10 if the state is invalid (e.g. a wall or out-of-bounds).
100         - -1 for any other (valid) step.
101         """
102         if state == self.end:
103             return 100
104         if not is_valid_state(state[0], state[1]):
105             return -10
106         return -1
107
108     # -----
109     # Q-learning Training Loop
110     # -----
111     for episode in range(num_episodes):
112         state = self.start
113         steps = 0
114         while steps < max_steps:
115             # Choose an action from the current state
116             action = choose_action(state)
117             x, y = state
118             dx, dy = actions[action]
119             new_state = (x + dx, y + dy)
120             reward = get_reward(new_state)
121
122             # Determine the maximum Q-value for the new state (if valid)
123             if is_valid_state(new_state[0], new_state[1]):
124                 best_next_q = np.max(q_table[new_state[0], new_state[1]])
125             else:
126                 best_next_q = 0
127
128             # Update the Q-value for the current state and action
129             q_table[x, y, action] += alpha * (reward + gamma * best_next_q - q_table
130             [x, y, action])
131
132             # Terminate the episode if:
133             # - The agent reaches the goal.
134             # - The agent takes an invalid move (hits a wall).
135             if new_state == self.end or not is_valid_state(new_state[0], new_state
136             [1]):
137                 break

```

```

135         # Otherwise, move to the new state and continue
136         state = new_state
137         steps += 1
138
139     # -----
140     # Extract the Optimal Path
141     # -----
142     path = [self.start]
143     current_state = self.start
144     visited = set([self.start]) # To help prevent loops
145
146     # Follow the greedy policy derived from the Q-table until the end is reached.
147     while current_state != self.end:
148         action = np.argmax(q_table[current_state[0], current_state[1]])
149         dx, dy = actions[action]
150         new_state = (current_state[0] + dx, current_state[1] + dy)
151
152         # If the new state is invalid or we've already visited it, break out to
153         # avoid an infinite loop.
154         if not is_valid_state(new_state[0], new_state[1]) or new_state in visited:
155             break
156
157         path.append(new_state)
158         visited.add(new_state)
159         current_state = new_state
160
161     return path
162
163 # Game flow
164 if __name__ == "__main__":
165     width, height = 21, 21
166
167     # Initialize simulator
168     simulator = MazeSimulator(width, height)
169     simulator.generate_maze()
170     simulator.set_start_and_end()
171
172     # Generate and display maze
173     maze_array = simulator.get_maze_array()
174     print("Maze array:")
175     for row in maze_array:
176         print(row)
177     print(f"Start point: {simulator.start}")
178     print(f"End point: {simulator.end}")
179
180     # Solve maze using RL
181     path = simulator.mazeSolverUsingRL() # Call the solver
182     print(f"\nRL Agent Path ({len(path)} steps):")
183     print(path)
184
185     # Save visualization with path
186     simulator.save_maze_image("maze.png")
187     print("\nMaze image saved as 'maze.png'")

```

Listing 1: Maze Simulator Code

5 Conclusion

The implemented maze simulator and Q-learning agent provide a functional solution for generating and navigating mazes. The agent learns to find paths through the maze, demonstrating the effectiveness of reinforcement learning techniques in solving pathfinding problems. The performance of the agent can be further improved by fine-tuning the hyperparameters and exploring other reinforcement learning algorithms. Future work could also involve visualizing the agent's learning process and comparing the performance of different maze-solving algorithms.